

Structural Admissibility for Verification Systems: Warrant Debt and Observational Quotients in Rocq

Duston Moore^{1*}

^{1*}Independent Researcher, Edmonton, Alberta, Canada.

Corresponding author(s). E-mail(s): dhwcmoore@gmail.com;

Abstract

Verification can fail before proof search begins. When the observation interface of a verification system conflates states that a safety predicate must distinguish, the verification problem is structurally malformed: the required information has already been destroyed by the observational quotient. We formalise this condition as structural admissibility. A predicate is admissible exactly when it factors through the observational quotient. Equivalently, the quotient must preserve every distinction required by the predicate. For perturbation systems, we introduce warrant debt as an operational witness of admissibility failure: warrant debt occurs when algebraic observation treats two perturbations as equivalent while empirical evidence distinguishes them. The Warrant Debt Theorem is proved in Rocq and shows that whenever warrant debt exists, the corresponding empirical tolerance predicate is not admissible. The mechanisation treats inadmissibility as a proof obligation prior to proof search. The Rocq development proves the Warrant Debt Theorem, the concrete rupture decomposition theorem for the `SimpleSynergyCompose` instance, and the Non-Factorisation Theorem. The abstract composition theorem is stated as an interface condition. The concrete instance satisfying that interface is proved. A higher-order dependency extension is presented as a mathematical generalisation, not as part of the present mechanisation.

Keywords: Formal verification, Safety-critical systems, System reliability, Observational quotients, Warrant debt, Rocq mechanisation

1 Introduction

The question is not whether a correctness predicate is true, but whether truth survives quotienting by the observation structure.

Verification failure is commonly diagnosed as proof failure: the prover did not find a certificate, the invariant was too weak, or the abstraction was too coarse. This diagnosis is often wrong.

Many failures are structural: the observation map has already collapsed the distinctions the correctness predicate requires before any proof attempt begins. In that situation, proof search is not merely difficult. The question it is attempting to answer is malformed.

In system design terms, structural admissibility is a design-time safety property. It asks whether the observation interface of a monitored or verified system preserves the distinctions required by the safety predicate. If the interface quotient collapses those distinctions, no downstream verifier, monitor, or proof-search procedure can recover them. The failure is therefore not merely a weakness of automation. It is a defect in the system's verification architecture.

When an observation map $M : X \rightarrow O$ is not injective, distinct states share an observational image. The induced equivalence

$$x \sim_M y \iff M(x) = M(y)$$

partitions the state space, and a predicate

$$\Phi : X \rightarrow \{0, 1\}$$

is *admissible* with respect to M if and only if it factors through that partition:

$$\Phi = \hat{\Phi} \circ M$$

for some

$$\hat{\Phi} : O \rightarrow \{0, 1\}.$$

If Φ is not admissible, there exist states x, y such that $M(x) = M(y)$ and $\Phi(x) \neq \Phi(y)$. No reasoning that operates within the observational regime can certify Φ , because the regime has erased the distinction Φ requires. Inadmissibility is structural impossibility, not epistemic limitation.

The problem belongs to quotient semantics, abstract interpretation, observational equivalence, and definability theory. The issue is not the expressive weakness of the prover; it is whether the observational quotient is fine enough to preserve the distinctions the predicate demands.

In perturbation systems, this failure appears as *warrant debt*: algebraic observation treats two perturbations as equivalent while empirical evidence distinguishes them. Warrant debt is therefore the concrete witness that admissibility has failed.

Dependent type theory makes admissibility constraints explicit as proof obligations prior to theorem proving. The question is not whether a proof can be found,

Table 1 Mechanisation status of the main claims

Claim	Rocq identifier	Status
Warrant debt implies non-admissibility	<code>warrant_debt_implies_non_admissible</code>	Machine-checked
Concrete rupture decomposition (<code>SimpleSynergyCompose</code>)	<code>rupture_decomposition_proved</code>	Machine-checked
Non-Factorisation	<code>no_alg_factorisation</code>	Machine-checked
Abstract composition interface (<code>COMPOSITION_SIG</code>)	<code>rupture_decomposition</code>	Axiomatic interface
Higher-order dependency lattice	—	Conjecture, not mechanised

but whether the correctness predicate is even well-formed relative to the observation structure. The Rocq development presented here makes this precise: warrant debt is represented as a formal object, and the theorem that its existence implies non-admissibility is fully proved.

Contributions.

The paper makes three contributions to the design of structurally sound verification systems. The Factorisation Theorem (Section 2) defines when a safety predicate is well formed relative to an observation interface. It gives three equivalent characterisations: factorisation through the observation map, constancy on equivalence classes, and partition refinement. This theorem is the criterion used by the rest of the paper. The Warrant Debt Theorem (Section 4) provides a mechanised witness of design-time inadmissibility: whenever empirical evidence distinguishes algebraically equivalent perturbations, the corresponding tolerance predicate cannot be verified within the current observational regime; the theorem is fully mechanised in Rocq. The mechanised development then separates two further issues. Section 6 proves a concrete binary rupture decomposition theorem for `SimpleSynergyCompose`. Section 7 proves the Non-Factorisation Theorem, which rules out total algebraic projections that preserve rupture behaviour in general. A higher-order extension to finite dependency families on n components is set out at the level of mathematical structure as a proposed generalisation; the higher-order theorem is not part of the present mechanisation.

2 Observational Quotients and Admissibility

This section states the abstract theorem of which all subsequent results are instances. No perturbation-specific machinery is required here.

2.1 Observation Maps and Observational Equivalence

Let X be a non-empty set (the state space) and O a non-empty set (the observation space). Let

$$M : X \rightarrow O$$

be a function, the *observation map*; M need not be injective.

Definition 2.1 (Observational Equivalence). The *observational equivalence* induced by M is

$$x \sim_M y \stackrel{\text{def}}{\iff} M(x) = M(y).$$

This is an equivalence relation on X .

Let $X/\sim_M$ denote the quotient and

$$\pi : X \rightarrow X/\sim_M$$

the canonical projection.

There exists a unique injection

$$\bar{M} : X/\sim_M \rightarrow O$$

such that

$$M = \bar{M} \circ \pi.$$

The observational equivalence is the relevant object of study. Two states that M cannot separate are, from the perspective of the verification system, identical. No predicate that respects the observation structure can distinguish them.

2.2 Admissible Predicates

Definition 2.2 (Admissibility). Let

$$\Phi : X \rightarrow \{0, 1\}$$

be a correctness predicate.

Φ is *admissible with respect to M* if and only if there exists

$$\tilde{\Phi} : X/\sim_M \rightarrow \{0, 1\}$$

such that

$$\Phi = \tilde{\Phi} \circ \pi.$$

That is, Φ must factor through the quotient induced by observational equivalence. Equivalently, since

$$M = \bar{M} \circ \pi,$$

admissibility is equivalent to factorisation through the observation map itself: there exists

$$\hat{\Phi} : O \rightarrow \{0, 1\}$$

such that

$$\Phi = \hat{\Phi} \circ M.$$

The quotient is primary. The codomain O matters only because the observation map induces the equivalence relation that defines the quotient. Verification is therefore governed not by the observation values themselves, but by the partition they impose on the state space.

2.3 The Factorisation Theorem

Theorem 2.3 (Factorisation). *The following conditions are equivalent.*

- (1) Φ is admissible with respect to M .
- (2) Φ is constant on $\sim_M$ -equivalence classes:

$$M(x) = M(y) \Rightarrow \Phi(x) = \Phi(y).$$

- (3) The observational equivalence refines the predicate equivalence:

$$\sim_M \subseteq \sim_\Phi,$$

where

$$x \sim_\Phi y \stackrel{\text{def}}{\iff} \Phi(x) = \Phi(y).$$

Proof (1) $\Rightarrow$ (2): If

$$\Phi = \hat{\Phi} \circ M$$

and

$$M(x) = M(y),$$

then

$$\Phi(x) = \hat{\Phi}(M(x)) = \hat{\Phi}(M(y)) = \Phi(y).$$

(2) $\Rightarrow$ (3): If

$$x \sim_M y,$$

then

$$M(x) = M(y),$$

so by (2),

$$\Phi(x) = \Phi(y),$$

hence

$$x \sim_\Phi y.$$

Therefore

$$\sim_M \subseteq \sim_\Phi.$$

(3) $\Rightarrow$ (1): Since

$$\sim_M \subseteq \sim_\Phi,$$

the assignment

$$\tilde{\Phi}([x]_M) = \Phi(x)$$

is well-defined on the quotient

$$X/\sim_M.$$

Thus

$$\Phi = \tilde{\Phi} \circ \pi.$$

Since $\bar{M}$ is injective into the image of M , this gives a factorisation through $\text{im}(M)$. A factorisation through all of O is obtained by assigning arbitrary Boolean values outside $\text{im}(M)$. $\square$

Theorem 2.3 is the central criterion used throughout the paper.

A predicate is decidable from observation if and only if the observation-induced partition is fine enough to respect the predicate's partition.

If the observational equivalence is too coarse, no argument about the predicate can recover what the quotient has already destroyed.

2.4 Saturation

Definition 2.4 (Saturation). A set

$$S \subseteq X$$

is *M-saturated* if for all

$$\begin{aligned} x \in S \quad \text{and} \quad y \in X, \\ M(x) = M(y) \Rightarrow y \in S. \end{aligned}$$

That is, S is a union of $\sim_M$ -equivalence classes.

Corollary 2.5 (Saturation Criterion). Φ is admissible with respect to M if and only if

$$\Phi^{-1}(1)$$

is *M-saturated*.

Proof Admissibility holds if and only if no $\sim_M$ -equivalence class straddles the boundary between

$$\Phi^{-1}(1) \quad \text{and} \quad \Phi^{-1}(0).$$

This holds if and only if

$$\Phi^{-1}(1)$$

is a union of $\sim_M$ -classes, which is exactly the definition of M -saturation. $\square$

2.5 Inadmissibility as Structural Impossibility

When Φ is not admissible, there exist

$$x, y \in X$$

such that

$$M(x) = M(y) \quad \text{and} \quad \Phi(x) \neq \Phi(y).$$

The pair (x, y) is an inadmissibility witness.

No predicate that factors through the current observation map can recover this distinction while that observational regime remains fixed. The quotient

$$X/\sim_M$$

has already collapsed the distinction that Φ requires.

Inadmissibility is not epistemic ignorance. It is structural impossibility.

The failure lies in the representational regime itself, not in the strength of the prover operating within it.

2.6 Literature Placement

Theorem 2.3 is a definability result. Its connections to the literature are structural.

In abstract interpretation [3, 4], a Galois connection (α, γ) relates concrete and abstract domains; a property survives abstraction only if the abstract domain is fine enough to preserve the required distinctions [5, 8]. The quotient $X/\sim_M$ is the abstract domain. Inadmissibility is the case in which the abstraction is too coarse and the required distinction has been irreversibly erased.

Beth’s theorem [1] establishes that implicit and explicit definability coincide in first-order logic. In the present setting there is a fixed observation map rather than a first-order theory, and the observation-induced equivalence plays the role of the theory’s indistinguishability relation. A predicate is definable in the relevant sense if and only if the quotient is fine enough to preserve its partition [9].

Two structures are elementarily equivalent if and only if the Duplicator wins the Ehrenfeucht–Fraïssé game of all finite lengths [6, 7]. The relation $\sim_M$ is the observational analogue: two states are $\sim_M$ -equivalent when no observation separates them. Inadmissibility is the corresponding structural result: collapse of the observational quotient forces indistinguishability between states that the correctness predicate must separate [10].

Milner’s observational equivalence [11, 12] identifies processes that no context can distinguish by observable actions. Observation maps play the role of contexts, and the inadmissibility condition is the verification analogue of contextual equivalence [13].

The relation to abstraction refinement in CEGAR-style verification [2] is also direct. In CEGAR, spurious counterexamples indicate that an abstraction is too coarse for the property being checked, and the repair is refinement of the abstraction after a failed proof attempt. Structural admissibility separates a prior failure mode: the observation quotient may already be too coarse for the predicate to be well-formed as a verification target. The corresponding repair is not refinement after failure but admissibility checking before proof search begins. Warrant debt makes the failure witness explicit rather than leaving it to be discovered by a spurious counterexample.

3 Concrete Perturbation Space

The abstract framework of Section 2 is instantiated here.

The observation map M becomes an algebraic observation map over perturbations: it preserves carrier class and magnitude while discarding emergent boundary information.

The correctness predicates Φ become tolerance predicates over interface regions.

This section introduces the formal objects required for the Warrant Debt Theorem of Section 4.

3.1 Boundary Type Classification

Definition 3.1 (Boundary Types). A perturbation is classified by one of three boundary types.

- (i) **INV** (invariant): tolerance conditions are fully determined at design time by algebraic structure;
- (ii) **REL** (relational): tolerance conditions vary within bounded and certifiable limits;
- (iii) **EMG** (emergent): tolerance conditions depend on information that cannot be recovered from the algebraic observation kernel alone.

The distinction is representational. **EMG** marks the case in which algebraic observation is structurally insufficient.

```

1 Inductive BType : Type :=
2   | INV  (* invariant: structurally determined *)
3   | REL  (* relational: bounded, variable *)
4   | EMG. (* emergent: observationally lost *)
5
6 Record Perturbation := mkPerturbation {
7   p_btype      : BType;
8   p_carrier    : CarrierClass;
9   p_magnitude : Interval;
10  p_mode       : nat
11 }.

```

Listing 1 Boundary types and perturbation record.

The **Interval** type carries a proof of well-formedness $\ell \leq h$, so every magnitude is represented as a certified interval rather than an uncertified real pair.

A structural consequence used later is that **EMG** is the top element under composition: if either component is emergent, the composition is emergent. Composition cannot eliminate emergent character.

3.2 Tolerance Predicates

Definition 3.2 (Tolerance Predicates). A tolerance predicate specifies whether an interface region can absorb a perturbation without violating its safety boundary. Two forms are distinguished: formal tolerance, decided at design time by a Boolean procedure; and empirical tolerance, indexed by evidence such as sealed monitoring logs or certified test records.

```

1 Definition TolPredicate := Perturbation -> Prop.

```



```

2
3 Definition formal_tol (dec : Perturbation -> bool)
4   : TolPredicate :=
5   fun p => dec p = true.
6
7 Parameter Evidence : Type.
8 Parameter empirical_tol : Evidence -> TolPredicate.

```

Listing 2 Tolerance predicates.

The type `Evidence` is left abstract. The results of Section 4 do not depend on the internal structure of evidence, only on the fact that empirical tolerance may distinguish perturbations that algebraic observation treats as equivalent.

This gap is the source of warrant debt.

3.3 Rupture

Definition 3.3 (Rupture). A perturbation p *ruptures* tolerance predicate `tol` when

$$\neg \text{tol}(p).$$

```

1 Definition Rupture
2   (tol : TolPredicate)
3   (p : Perturbation) : Prop :=
4   ~ tol p.

```

Listing 3 Rupture.

Rupture is not a property of the perturbation alone. It is a relation between a perturbation and a tolerance predicate.

Changing the predicate changes what counts as rupture.

This matters because detectability depends not only on whether rupture occurs, but on whether the tolerance predicate itself is admissible with respect to the observation structure. Section 4 makes this precise.

4 Algebraic Observation and Warrant Debt

This section is the formal centre of the paper.

Section 2 established the abstract admissibility criterion: a predicate is admissible if and only if the observational equivalence induced by the observation map is fine enough to preserve the distinctions the predicate requires.

This section is the concrete instance of that criterion for perturbation systems.

The algebraic observation map plays the role of M ; `alg_equiv` plays the role of $\sim_M$; and `PredAdmissible` is the perturbation-level factorisation condition of Theorem 2.3.

4.1 The Algebraic Observation Map

The algebraic observation map strips emergent boundary information.

The relation `alg_equiv` is an equivalence relation on the algebraically observable subspace of perturbations. On the full `Perturbation` type it is a partial equivalence

relation. Emergent perturbations lie outside its domain: since `btype_algebraic(EMG) = false`, no emergent perturbation is related even to itself.

```

1 Definition btype_algebraic (b : BType) : bool :=
2   match b with
3   | INV => true
4   | REL => true
5   | EMG => false
6   end.
7
8 Definition alg_equiv (p q : Perturbation) : Prop :=
9   p_carrier p = p_carrier q /\
10  p_magnitude p = p_magnitude q /\
11  btype_algebraic (p_btype p) = true /\
12  btype_algebraic (p_btype q) = true.
13
14 Definition PredAdmissible
15   (P : Perturbation -> Prop) : Prop :=
16   forall p q : Perturbation,
17     alg_equiv p q -> (P p <-> P q).

```

Listing 4 Algebraic equivalence and predicate admissibility.

Definition 4.1 (Observational Admissibility in Perturbation Space). A predicate P is *admissible* with respect to the algebraic observation map if and only if it is invariant under `alg_equiv`: P cannot distinguish perturbations that the algebraic kernel cannot distinguish.

This is the perturbation-level realisation of condition (2) of Theorem 2.3:

$$M(x) = M(y) \Rightarrow \Phi(x) = \Phi(y).$$

Emergent perturbations lie outside the algebraic observation domain entirely:

$$\text{btype_algebraic}(\text{p_btype } q) = \text{false}$$

whenever

$$\text{p_btype } q = \text{EMG}.$$

The algebraic kernel therefore forms no equivalence class containing an emergent perturbation. Predicates separating algebraic and emergent perturbations are not ruled out by `PredAdmissible`; rather, they fall outside the algebraic observation regime. This is why the Non-Factorisation Theorem in Section 7 is needed: it shows that no total projection from the full perturbation space into the algebraic fragment can preserve rupture behaviour for all tolerance predicates.

4.2 Warrant Equivalence

Definition 4.2 (Warrant Equivalence). Perturbations p and q are *warrant-equivalent* with respect to evidence e when empirical tolerance treats them identically:

$$\text{warrant_equiv}(e, p, q) \stackrel{\text{def}}{\iff} \left(\text{empirical_tol}(e, p) \leftrightarrow \text{empirical_tol}(e, q) \right).$$

```

1 Definition warrant_equiv
2   (e : Evidence)
3   (p q : Perturbation) : Prop :=
4   empirical_tol e p <-> empirical_tol e q.

```

Listing 5 Warrant equivalence.

Warrant equivalence is the empirical counterpart of algebraic equivalence.

Algebraic equivalence records what the formal certification infrastructure treats as the same perturbation.

Warrant equivalence records what the empirical evidence treats as the same perturbation.

Warrant debt arises precisely when these two notions of sameness disagree.

4.3 Warrant Debt

Definition 4.3 (Warrant Debt). A *warrant debt* exists between perturbations p and q with respect to evidence e when they are algebraically equivalent but not warrant-equivalent:

$$\text{WarrantDebt}(e, p, q) \stackrel{\text{def}}{\iff} \text{alg_equiv}(p, q) \wedge \neg \text{warrant_equiv}(e, p, q).$$

```

1 Definition WarrantDebt
2   (e : Evidence)
3   (p q : Perturbation) : Prop :=
4   alg_equiv p q /\ ~ warrant_equiv e p q.

```

Listing 6 Warrant debt.

A $\text{WarrantDebt}(e, p, q)$ is the formal witness that the observational regime is too coarse.

The algebraic certification infrastructure treats p and q as equivalent, while the evidence record distinguishes them.

Admissibility therefore fails at the level of representation itself: the system is attempting to certify a predicate that its own observation structure cannot support.

Example 4.4 (Warrant debt from collapsed observation). Let p and q be perturbations with the same carrier class and certified magnitude interval, both classified as algebraic by `btype_algebraic`. Then `alg_equiv p q` holds. If a sealed evidence record e certifies `empirical_tol e p` but not `empirical_tol e q`, then p and q are separated by the evidence while the algebraic observation regime identifies them. This is precisely $\text{WarrantDebt}(e, p, q)$, and Theorem 4.5 implies that `empirical_tol e` is not admissible under the current algebraic observation map. The correct response is refinement of the observation map, not renewed proof search inside the same quotient.

4.4 The Warrant Debt Theorem

Theorem 4.5 (Warrant Debt Implies Non-Admissibility). *For any evidence e and perturbations p and q ,*

$$\text{WarrantDebt}(e, p, q) \Rightarrow \neg \text{PredAdmissible}(\text{empirical_tol}(e)).$$

Proof Suppose

$$\text{WarrantDebt}(e, p, q).$$

Then

$$\text{alg_equiv}(p, q)$$

and

$$\neg \text{warrant_equiv}(e, p, q).$$

Assume for contradiction that

$$\text{PredAdmissible}(\text{empirical_tol}(e)).$$

By Definition 4.1, admissibility requires that algebraically equivalent perturbations cannot be separated by the predicate.

Hence

$$\text{alg_equiv}(p, q)$$

implies

$$\text{empirical_tol}(e, p) \leftrightarrow \text{empirical_tol}(e, q),$$

which is exactly

$$\text{warrant_equiv}(e, p, q).$$

This is the contradiction. Admissibility forces warrant equivalence, while warrant debt asserts its negation. Both cannot hold simultaneously. $\square$

The Rocq mechanisation is fully proved:

```

1 Theorem warrant_debt_implies_non_admissible :
2   forall (e : Evidence)
3     (p q : Perturbation),
4     WarrantDebt e p q ->
5     ~ PredAdmissible (empirical_tol e).
6 Proof.
7   intros e p q [Hequiv Hne_warrant] Hadm.
8   apply Hne_warrant.
9   unfold PredAdmissible in Hadm.
10  exact (Hadm p q Hequiv).
11 Qed.

```

Listing 7 Warrant Debt Theorem, fully proved in Rocq.

The proof is short because the admissibility obstruction has been placed in the definitions rather than deferred to proof search.

The engineering consequence is immediate. If a warrant debt can be exhibited, the system is operating with an inadmissible observational regime. More proof cannot repair representational failure. The correct response is not stronger proof search, but refinement of the observation map until admissibility is restored, or formal recognition that the target predicate is not verifiable within the current regime.

Connection to Theorem 2.3.

The Warrant Debt Theorem is the concrete instance of the Factorisation Theorem for perturbation systems.

`alg_equiv` is the induced equivalence $\sim_M$.

`PredAdmissible` is the factorisation condition.

The pair (p, q) with `WarrantDebt`(e, p, q) is the inadmissibility witness of Theorem 2.3: two states the observation map conflates that the predicate must separate.

The implication of Theorem 4.5 is the perturbation-level form of condition (2) in Theorem 2.3.

Thus Theorem 4.5 is the perturbation-level instance of Theorem 2.3.

5 Type-Level Realisation

This section shows how admissibility appears as a proof obligation inside dependent type theory.

The type structure does not replace the mathematics; it exhibits the same obligations in machine-checkable form. Refinement of observation corresponds to refinement of the formal objects through which verification proceeds.

Theorem 4.5 established that warrant debt implies non-admissibility. The present section states the operational consequence: before proving a predicate, one must first establish that the predicate is well-formed relative to the available observation structure.

5.1 Observational Refinement as Product Extension

When a predicate Φ is inadmissible with respect to an observation map M_i , the corrective is not stronger proof search but observational refinement: introduce an auxiliary observation that separates the inadmissibility witnesses.

Definition 5.1 (Observational Refinement). Let

$$M_i : X \rightarrow O_i.$$

An *observational refinement* of M_i is a map

$$M_{i+1} : X \rightarrow O_{i+1}$$

such that

$$x \sim_{M_{i+1}} y \Rightarrow x \sim_{M_i} y.$$

That is, the refined observation induces a strictly finer partition. A common constructive form of refinement is product extension:

$$M_{i+1} = (M_i, N_i)$$

for some auxiliary observation

$$N_i : X \rightarrow F_i,$$

so that

$$O_{i+1} \cong O_i \times F_i.$$

Proposition 5.2 (Product Extension). *If*

$$M_{i+1} = (M_i, N_i)$$

is an observational refinement, then

$$O_{i+1} \cong O_i \times F_i,$$

and the projection

$$\pi_1 : O_{i+1} \rightarrow O_i$$

satisfies

$$M_i = \pi_1 \circ M_{i+1}.$$

The refined observation reduces to the original by projection.

Each refinement step therefore extends the type of the observation space by a product factor. The projection π_1 corresponds exactly to forgetting the new factor.

Refinement is not replacement. It is conservative extension: the old observation remains recoverable by projection, while the new factor provides the separation required to eliminate inadmissibility.

5.2 Type Structure and Proof Obligations

Type structure enforces the projection and refinement obligations corresponding to Theorem 2.3.

The issue is not whether a proof can be found, but whether the predicate is even well-formed relative to the observation structure.

A predicate admissible with respect to M_i can be stated without appeal to distinctions outside the partition induced by M_i . A non-admissible predicate cannot: any attempt to verify it requires appealing to a finer partition, which is exactly the formal signal that the observation map must be refined.

In the Rocq development, $\text{PredAdmissible}(P)$ is the explicit proof obligation that must be discharged before verification of a tolerance predicate may proceed. If it cannot be discharged, the observation map must be refined by product extension as in Proposition 5.2, or it must be formally certified that P is not verifiable within the current regime. Only after admissibility is established does proof search for P begin.

This ordering is not a recommendation about proof strategy. It is a structural consequence of Theorem 4.5.

Warrant debt proves that proof search inside an inadmissible observational regime is not incomplete but misdirected: the system is attempting to certify a predicate its own representation cannot support.

6 Composition and Rupture Decomposition

6.1 Composition Signature

Composition is expressed as a Rocq module type. A compliant operator must satisfy three axioms: emergent type propagation, magnitude monotonicity, and a bridge between composite rupture and its sources.

```

1 Module Type COMPOSITION_SIG.
2   Parameter par_compose :
3     Perturbation -> Perturbation -> Perturbation.
4
5   Parameter synergy_debt :
6     Perturbation -> Perturbation -> Prop.
7
8   Axiom compose_btype_emg :
9     forall p q,
10      p_btype p = EMG /\ p_btype q = EMG ->
11      p_btype (par_compose p q) = EMG.
12
13   Axiom compose_magnitude_ub :
14     forall p q,
15      (iv_hi (p_magnitude p) <= iv_hi (p_magnitude (par_compose p q)))%R /\
16      (iv_hi (p_magnitude q) <= iv_hi (p_magnitude (par_compose p q)))%R.
17
18   Axiom rupture_synergy_bridge :
19     forall (tol : TolPredicate) (p q : Perturbation),
20      Rupture tol (par_compose p q) ->
21      Rupture tol p /\ Rupture tol q /\ synergy_debt p q.
22 End COMPOSITION_SIG.

```

Listing 8 Abstract composition module signature.

Definition 6.1 (Interaction Debt, Binary Case). Two perturbations incur *interaction debt* (in this binary instance also called *synergy debt*) when both are individually tolerated under a predicate, yet their composition ruptures under that same predicate. The binary case names the lowest non-trivial level of the interaction debt hierarchy and is not reducible to component-level failure. The general hierarchy, which classifies interaction debt by the size of the minimal support set on which inadmissibility arises, is developed in Section 6.3.

6.2 Proved Rupture Decomposition

The concrete module `SimpleSynergyCompose` satisfies the axioms of `COMPOSITION_SIG` with full proofs. The rupture decomposition is therefore established constructively.

Theorem 6.2 (Rupture Decomposition). *For any tolerance predicate `tol` and perturbations p, q ,*

$$\text{Rup}(\text{tol}, \text{par_compose}(p, q)) \Rightarrow \text{Rup}(\text{tol}, p) \vee \text{Rup}(\text{tol}, q) \vee \text{synergy_debt}(p, q).$$

```

1 Lemma rupture_synergy_bridge :

```

```

2 forall (tol : TolPredicate) (p q : Perturbation),
3   Rupture tol (par_compose p q) ->
4   Rupture tol p \/\ Rupture tol q \/\ synergy_debt p q.
5 Proof.
6   intros tol p q Hrup.
7   destruct (classic (tol p)) as [Hp | Hnp].
8   - destruct (classic (tol q)) as [Hq | Hnq].
9     + right. right. unfold synergy_debt. exists tol.
10      exact (conj Hp (conj Hq Hrup)).
11      + right. left. exact Hnq.
12    - left. exact Hnp.
13 Qed.
14
15 Theorem rupture_decomposition_proved :
16   forall (tol : TolPredicate) (p q : Perturbation),
17   Rupture tol (SimpleSynergyCompose.par_compose p q) ->
18   Rupture tol p \/\ Rupture tol q \/\
19   (exists tol' : TolPredicate,
20    tol' p /\ tol' q /\
21    Rupture tol' (SimpleSynergyCompose.par_compose p q)).
22 Proof.
23   intros tol p q Hrup.
24   exact (SimpleSynergyCompose.rupture_synergy_bridge tol p q Hrup).
25 Qed.

```

Listing 9 Proved rupture decomposition.

The third disjunct identifies the interaction term: both components satisfy a tolerance predicate, yet their composition violates it. Interaction debt at the binary level is therefore not residual; it is a named and typed obligation.

Remark 6.3 (Classical reasoning). The proof of `rupture_synergy_bridge` uses classical case analysis over `tol p` and `tol q` via the `classic` tactic, which invokes excluded middle for arbitrary `TolPredicate : Perturbation -> Prop`. This is appropriate for a result stated in `Prop`. An executable monitoring instance would require a decidable tolerance predicate, such as `formal_tol dec`, and the corresponding decomposition should be stated with an explicit decidability assumption.

The role of Theorem 6.2 is compositional. Where Theorem 4.5 identifies failure of admissibility at the level of a single predicate, Theorem 6.2 identifies the sources of failure in composed systems for the binary case. A composite rupture on two components decomposes into certified local failures or a single binary interaction term. The binary formulation is, however, only the first non-trivial case of a hierarchy. For systems with three or more components, a single residual interaction term is no longer sufficient: it records that some debt exists without identifying the minimal locus of failure. Section 6.3 develops the corresponding hierarchy.

6.3 Higher-Order Generalisation

The following construction is not part of the present Rocq mechanisation. It identifies the finite dependency-lattice structure that a future mechanisation must certify.

The binary formulation of rupture decomposition is only the first non-trivial case. In general, for a system with component family $C = \{C_1, \dots, C_n\}$, global inadmissibility need not reduce to one residual interaction term. Such a residual merely records that some interaction debt exists; it does not identify the minimal locus of failure.

The generalisation is hierarchical. For each non-empty subset $S \subseteq C$, let Σ_S denote the warrant debt associated with the interaction structure of the components in S , namely the failure of the global predicate to factor through the observation map restricted to that subsystem. Then local failure is the case $|S| = 1$, pairwise interaction debt is the case $|S| = 2$ recovered by Theorem 6.2, triple interaction debt the case $|S| = 3$, and so forth.

Accordingly, global rupture can be analysed through the disjunctive family of support witnesses

$$\text{Rup}(\Phi_{\text{global}}) \Rightarrow \bigvee_{\emptyset \neq S \subseteq C} \Sigma_S,$$

where Σ_S records that S supports a witness of inadmissibility. The question is not whether some Σ_S appears. The question is which S is minimal.

Definition 6.4 (Dependency Family and Minimal Support). The *dependency family* of a global predicate Φ over the component family C is

$$\mathcal{D}(\Phi) = \{ S \subseteq C \mid \Phi \text{ requires } S \text{ for admissibility} \}.$$

A subset $S \in \mathcal{D}(\Phi)$ is a *minimal support set* when no proper subset of S lies in $\mathcal{D}(\Phi)$. The minimal elements of $\mathcal{D}(\Phi)$ identify the smallest subsystems on which Φ genuinely depends. This definition is schematic: a full formalisation requires explicit projection maps M_S and projected predicates $\Phi|_S$ for each subsystem $S \subseteq C$.

A predicate may be admissible on every proper projection of size less than k , yet fail to be admissible on the full system. This is k -th order synergy.

Conjecture 6.5 (Higher-Order Warrant Debt). Let Φ be a global predicate on a finite component family $C = \{C_1, \dots, C_n\}$, with admissibility-induced dependency family $\mathcal{D}(\Phi)$. Suppose that Φ is inadmissible relative to the observation map on C . If, for some $k \leq n$, every projected predicate $\Phi|_T$ is admissible for every subsystem $T \subseteq C$ with $|T| < k$, then there exists a minimal support set $S^* \in \mathcal{D}(\Phi)$ with $|S^*| \geq k$. The set S^* is the smallest subsystem on which higher-order warrant debt arises.

The conjecture rests on the powerset-lattice structure of the upward-closed family of inadmissibility witnesses. A proof-sketch argument proceeds by projecting the inadmissibility pair from the Factorisation Theorem onto every subset $S \subseteq C$ and taking the minimal element of the resulting family, which has cardinality at least k by hypothesis. The binary case ($k = 2$) is the content of Theorem 6.2 and is machine-checked in Rocq via `rupture_synergy_bridge` for `SimpleSynergyCompose`. A Rocq mechanisation of the general case would require a powerset-lattice development parameterised over $|C|$; this is identified as a proposed extension to the present development.

6.4 Industrial Safety-Lifecycle Interpretation

The industrial significance is practical as well as formal. In a safety-critical system, the observation map is not an incidental logging choice. It is part of the safety case.

Table 2 Safety-lifecycle role of structural admissibility

Lifecycle stage	Admissibility question	Action on failure
Design time	Does the observation map preserve the distinctions required by the safety predicate?	Refine sensor placement, state representation, or monitor interface
Deployment	Are monitor predicates admissible under the logged observation regime?	Reject monitor design or add evidence channel
Runtime	Does current evidence create warrant debt?	Trigger structural review, not ordinary anomaly triage
Post-failure audit	What is the minimal support set of the violated predicate?	Identify the subsystem, interface, or design assumption to be reviewed

Sensor placement, sampling policy, aggregation logic, alarm thresholds, and monitor interfaces jointly determine which distinctions are preserved for verification. Structural admissibility therefore belongs at design time, before monitor deployment and before proof obligations are discharged.

Consider an energy or semiconductor process in which pressure, temperature, compressor lag, flow timing, chamber state, or tool-cycle phase are each locally within tolerance. A local anomaly detector may classify every channel as benign, and a pairwise monitor may find no violation. Yet the joint configuration may violate a global invariant. The admissibility failure is not that the data are noisy. It is that the observational quotient supplied to the verifier has collapsed the distinction on which the global safety predicate depends. The claim is not that the present paper models those domains in detail, but that the admissibility obligation identifies a design condition any such monitor must satisfy.

The Rocq-certified warrant-debt witness makes this failure auditable. It identifies that algebraic equivalence and empirical evidence disagree, and therefore that the current observation interface is not adequate for the target safety predicate. In an industrial safety lifecycle, the witness has three uses. Before deployment, it can reject an insufficient monitor design. During operation, it can trigger structural review rather than ordinary anomaly triage. After failure, it can identify which observational boundary must be refined.

6.5 Bridge Theorem and Audit Interpretation

The bridge theorem is the engineering conversion from failure detection to post-failure audit: it turns a rupture event into an assignable minimal support set.

The bridge from rupture to audit therefore requires classification rather than mere existential detection. The bridge result must identify not only that interaction debt exists, but the minimal support set on which it arises. Only then can verification assign warrant debt: the failed observational boundary, the required refinement, and the subsystem, interface, or design assumption to be reviewed.

Remark 6.6 (Strengthened Rupture-to-Synergy Bridge). Assuming the higher-order dependency-family structure of Conjecture 6.5, a global rupture whose lower-order projections do not suffice determines a minimal support set $S^* \in \mathcal{D}(\Phi)$ of order at least k that witnesses higher-order warrant debt.

This is stronger than the binary bridge axiom of Section 6, which records existential interaction debt without classifying its order. The strengthened formulation replaces an unclassified residual with a formal object: the minimal witness of inadmissibility, located at a definite level of the dependency hierarchy.

The two mechanised results, Theorem 4.5 and Theorem 6.2, separate representational failure from binary compositional failure. Conjecture 6.5 indicates how this separation should extend to higher-order dependency families, but that extension is not claimed as a mechanised result of the present paper.

7 Mechanisation Boundary

The principal non-factorisation result is now fully proved in the Rocq development. In particular, the theorem `no_alg_factorisation` is established in the canonical development file `bdgi_perturbation_proved.v`. It proves, by explicit witness construction, that rupture cannot be uniformly preserved by projection into the algebraic fragment. This removes the earlier conjectural status of non-factorisation. It also makes the limitation structural, not epistemic.

The theorem should be read as a maximal-expressivity limit. It does not say that no useful algebraic abstraction exists. It says that once tolerance predicates are allowed to range over the full perturbation space, no total projection into the algebraic fragment can preserve rupture uniformly. Any useful algebraic verifier must therefore restrict the predicate class, refine the observation map, or explicitly certify the loss.

The theorem is therefore an impossibility result for verification regimes that attempt to certify rupture behaviour after projecting the perturbation space onto its algebraically observable fragment.

Theorem 7.1 (Non-Factorisation). *There is no algebraic homomorphism*

$$f : \text{Perturbation} \rightarrow \text{AlgPerturbation}$$

that preserves rupture for every tolerance predicate.

Equivalently, no total projection from the full perturbation space to its algebraically observable fragment can preserve rupture behaviour in general.

```

1 Definition AlgPerturbation :=
2   { p : Perturbation |
3     btype_algebraic (p_btype p) = true }.
4
5 Definition AlgHom :=
6   Perturbation -> AlgPerturbation.
7
8 Definition proj_alg (ap : AlgPerturbation) : Perturbation :=
9   proj1_sig ap.
10
11 Definition RupturePreserving

```

```

12 (f : AlgHom)
13 (tol : TolPredicate) : Prop :=
14   forall p : Perturbation,
15     Rupture tol p <->
16     Rupture tol (proj_alg (f p)).

```

Listing 10 Non-factorisation type signature.

The Rocq proof fixes an arbitrary $f : \text{AlgHom}$ and an emergent perturbation p_0 with $\text{p_btype}(p_0) = \text{EMG}$, then constructs a witness tolerance predicate that holds on every algebraic perturbation and fails at p_0 . This contradicts rupture preservation and discharges the proof obligation explicitly within the canonical development file. Thus the obstruction is not an implementation limitation of the chosen Rocq development. It is a limit theorem: any total algebraic projection of this form destroys rupture-relevant structure for some tolerance predicate.

Remark 7.2 (Correction of an earlier formulation). An earlier draft of this development presented a related claim, labelled `emg_sensitive_not_admissible`, asserting that any predicate sensitive to the emergent–algebraic distinction is inadmissible with respect to the algebraic observation map. Under the present definition of `alg_equiv` that formulation is not valid, and the corresponding lemma has been removed from the development rather than retained as an `Admitted` declaration. The witness-based non-factorisation result stated above is the correct theorem. It captures the substantive content of the earlier claim, namely that the algebraic regime cannot retain the distinctions on which rupture depends, without the invalid universal quantification over admissibility predicates present in the earlier formulation.

7.1 The Abstract Composition Boundary

The development contains an abstract interface statement `rupture_decomposition` over the opaque parameter `par_compose`. This statement is not counted among the proved results claimed in this paper. Since `par_compose` is introduced as an opaque composition operator, the statement is not derivable without concrete composition axioms, and no claim of its proof is made here. The proved result is the concrete theorem `rupture_decomposition_proved` for the explicit `SimpleSynergyCompose` instance, where the required bridge laws are established with full machine-checked proofs.

This distinction is methodological rather than provisional. The abstract statement specifies the interface conditions required of any admissible composition operator, and the concrete theorem certifies one realised instance satisfying those conditions. Any further module bound by `COMPOSITION_SIG` obtains the theorem at instantiation time, conditional on supplying the axioms the signature names.

7.2 Higher-Order Generalisation

Higher-order warrant debt for systems with more than two interacting components is not part of the present mechanisation. The current Rocq development proves the binary case, in which interaction debt first appears non-trivially. The finite dependency lattice required for the general case is mathematically straightforward but proof-theoretically lengthy, and is reserved for subsequent work.

8 Conclusion

For verification systems mediated by an observation interface, admissibility is prior to proof. Only after the target predicate is shown to respect the observational quotient does proof search address the predicate itself.

The Factorisation Theorem states the abstract criterion: a predicate is decidable from observation if and only if the observation-induced partition is fine enough to preserve the distinctions the predicate requires.

The Warrant Debt Theorem is the concrete operational instance of that criterion. When algebraically equivalent perturbations are distinguished by empirical evidence, the corresponding empirical tolerance predicate is not admissible. This result is fully mechanised in Rocq.

The Rupture Decomposition Theorem extends the account to composed systems. At the binary level, composite failure reduces to one of two cases: local component rupture or a named interaction term. The binary result is fully proved in Rocq for the concrete `SimpleSynergyCompose` instance, and the abstract theorem over `COMPOSITION_SIG` is provable for any module that supplies the required composition axioms.

The higher-order extension is not mechanised here. It indicates how the binary result should generalise to dependency families. In that setting, a global failure is classified by a minimal support set. That set identifies the smallest subsystem on which the global predicate depends, and hence the smallest subsystem, interface, or design assumption requiring review.

The Non-Factorisation Theorem is fully proved in the Rocq development and rules out any total algebraic projection that preserves rupture behaviour for every tolerance predicate.

Together, these results separate two distinct sources of verification failure. Representational failure occurs when the observational regime is too coarse and admissibility fails; it requires refinement of the observation map. Compositional failure occurs when admissibility holds but interaction produces rupture; it is analysed within the fixed observational regime.

More proof cannot repair the first. More abstraction cannot eliminate the second.

The connections to abstract interpretation [3, 4], Ehrenfeucht–Fraïssé indistinguishability [6, 10], Beth definability [1, 9], and observational equivalence [11, 12] are structural.

They place admissibility inside an established family of results concerning abstraction, quotienting, and definability: when the quotient destroys the distinction, proof cannot restore it.

The type `WarrantDebt`(e, p, q) is the formal witness that the empirical record preserves a distinction the algebraic regime has excluded.

Structural admissibility is therefore a design-time condition on verification systems. A monitor or prover is not merely asked to certify a predicate. It must first expose enough structure for that predicate to be meaningful under the system’s observation regime.

Statements and Declarations

Funding.

The author received no funding, grants, or other financial support for the research, authorship, or preparation of this manuscript.

Competing interests.

The author has no competing financial or non-financial interests to disclose that are relevant to the content of this article.

Data availability.

This is a formal verification paper and contains no empirical datasets. The Rocq source files corresponding to the mechanised results reported in Sections 4–7, including the proofs of the Warrant Debt Theorem, the concrete rupture decomposition theorem for `SimpleSynergyCompose`, and the Non-Factorisation Theorem, are available in the canonical development file `bdgi_perturbation_proved.v` at <https://github.com/dhwcmoore/structural-admissibility-in-rocq>.

Code availability.

The Rocq development is distributed under an open licence at the repository cited above; the exact tested version and build command are specified in the repository README.

Author contributions.

Duston Moore is the sole author and was responsible for the conception, formal development, Rocq mechanisation, manuscript preparation, and final approval of the submitted version.

Ethics approval.

Not applicable. The research did not involve human participants, human tissue, or animals.

Use of AI-assisted tools.

AI-assisted tools were used for copy-editing and editorial support. The author reviewed and accepts responsibility for the final manuscript.

References

- [1] E. W. Beth. On Padoa’s method in the theory of definition. *Indagationes Mathematicae*, 15:330–339, 1953.
- [2] E. M. Clarke, O. Grumberg, S. Jha, Y. Lu, and H. Veith. Counterexample-guided abstraction refinement for symbolic model checking. *Journal of the ACM*, 50(5):752–794, 2003.

- [3] P. Cousot and R. Cousot. Abstract interpretation: A unified lattice model for static analysis of programs by construction or approximation of fixpoints. In *Proceedings of the 4th ACM Symposium on Principles of Programming Languages*, pages 238–252. ACM, 1977.
- [4] P. Cousot and R. Cousot. Systematic design of program analysis frameworks. In *Proceedings of the 6th ACM Symposium on Principles of Programming Languages*, pages 269–282. ACM, 1979.
- [5] P. Cousot. Constructive design of a hierarchy of semantics of a transition system by abstract interpretation. *Theoretical Computer Science*, 277(1–2):47–103, 2002.
- [6] A. Ehrenfeucht. An application of games to the completeness problem for formalized theories. *Fundamenta Mathematicae*, 49:129–141, 1961.
- [7] R. Fraïssé. Sur quelques classifications des systèmes de relations. *Publications Scientifiques de l’Université d’Alger*, 1:35–182, 1954.
- [8] R. Giacobazzi, F. Ranzato, and F. Scozzari. Making abstract interpretations complete. *Journal of the ACM*, 47(2):361–416, 1998.
- [9] W. Hodges. *Model Theory*. Cambridge University Press, 1993.
- [10] L. Libkin. *Elements of Finite Model Theory*. Springer, 2004.
- [11] R. Milner. A calculus of communicating systems. *Lecture Notes in Computer Science*, 92, 1980.
- [12] R. Milner. *Communication and Concurrency*. Prentice Hall, 1989.
- [13] D. Sangiorgi. *Introduction to Bisimulation and Coinduction*. Cambridge University Press, 2011.